

Real Time Systems: Scheduling Algorithms

Index

- [Introduction](#)
- [Concepts](#)
 - [Hyperperiod](#)
 - [Secondary Period](#)
 - [Secondary Period: First Condition](#)
 - [Secondary Period: Second Condition](#)
 - [Secondary Period: Third Condition](#)
 - [Response Time Analysis](#)
 - [Critical Time](#)
 - [Absolute Deadline](#)
 - [Contribution](#)
 - [Processor on Demand Criterion](#)
 - [Arrival Time](#)
 - [Aperiodic tasks](#)
- [Cyclic Scheduler](#)
 - [Cyclic Scheduler: Requisites](#)
 - [Cyclic Scheduler: Methodology](#)
 - [Cyclic Scheduler: Pros](#)
 - [Cyclic Scheduler: Cons](#)
- [Rate Monotonic](#)
 - [Rate Monotonic: Requisites](#)
 - [Rate Monotonic: Methodology](#)
 - [Rate Monotonic: Methodology > Condition 1](#)
 - [Rate Monotonic: Methodology > Condition 2](#)
 - [Rate Monotonic: Pros](#)
 - [Rate Monotonic: Cons](#)
- [Deadline Monotonic](#)
 - [Deadline Monotonic: Requisites](#)
 - [Deadline Monotonic: Methodology](#)
 - [Deadline Monotonic: Pros](#)
 - [Deadline Monotonic: Cons](#)
- [Earliest Deadline First](#)
 - [Earliest Deadline First: Requisites](#)
 - [Earliest Deadline First: Methodology](#)
 - [Earliest Deadline First: Methodology > Processor on Demand Criterion](#)
 - [Earliest Deadline First: Precedence](#)
 - [Earliest Deadline First: New Task](#)
 - [Earliest Deadline First: Pros](#)
 - [Earliest Deadline First: Cons](#)

Introduction

Various basic scheduling algorithms. For these to work, we need more information about the task other than it's [basics](#). Now each task also has a priority P .

Concepts

New concepts

Hyperperiod

A hyperperiod H is period where the set of tasks will go back to the base line. It is computed as the *least common factor* of all [periods](#). Formally written:

$$H = \text{lcm}_{1 \leq i \leq n} T_i$$

Secondary Period

A secondary period where it is fitted between the longest computing time, shortest deadline. It has 2 other conditions. Formally defined:

$$\left\{ \begin{array}{l} T_s \in [\min_{1 \leq i \leq n} D_i, \max_{1 \leq i \leq n} C_i] \\ H = kT_s, k \in \mathbb{Z}^+ \\ \forall i : 2T_s - \gcd(T_s, T_i) \leq D_i \end{array} \right.$$

Secondary Period: First Condition

The first condition of the secondary period is:

$$T_s \in [\min_{1 \leq i \leq n} D_i, \max_{1 \leq i \leq n} C_i]$$

Basically it states that $T_{\{s\}}$ must be *grater than or equal* to the *maximum compute time* of all tasks. It also needs to be *less than or equal* the *shortest deadline* of all tasks.

Secondary Period: Second Condition

This condition is defined as:

$$H = kT_s, k \in \mathbb{Z}^+$$

This condition states that T_s is proportional to the hyperperiod H with a factor of k , with k being a positive complex (\mathbb{Z}^+) value.

Secondary Period: Third Condition

This condition is defined as:

$$\forall i : 2T_s - \gcd(T_s, T_i) \leq D_i$$

It basically means that for all tasks, the operation:

$$2T_s - \gcd(T_s, T_i)$$

Must be *less than or equal* to the deadline of the each task. This condition might fail and you can try *task splitting*. This technique consists on splitting a single task into different parts, always without the *critical section* in the edge of a section.

T_a is split into 2 parts :

$$\begin{cases} T_{a1} &= T_a \\ T_{a2} &= T_a \end{cases}$$
$$\begin{cases} D_{a1} &= D_a \\ D_{a2} &= D_a \end{cases}$$
$$\begin{cases} C_{a1} &= C_a \\ C_{a2} &= C_a \end{cases}$$

Response Time Analysis

This analysis checks the interference due to **preemption** of a higher priority task. Note that this only applies if $D_i \leq T_i$.

The response time R_i is based on the computing time c_i and the interference I_i produced by a higher priority task, preempting this lower priority one. It is defined as:

$$\begin{aligned}\forall \tau_i : R_i &= C_i + I_i \\ I_i &= \sum_{j \in hp(i)} \lceil \frac{R_i}{T_j} \rceil C_j \\ hp(i) &= \{j : 1..n | P_j > P_i\}\end{aligned}$$

The function $hp(i)$ start for *higher priority* and it is a set of tasks that have higher priority than the current one.

The solution basically works using this function:

$$W_i^{n+1} = c_i + \sum_{j \in hp(i)} \lceil \frac{W_i^n}{T_j} \rceil C_j$$

With $W_i^0 = c_i$, corresponding to the response time of task i which cannot be preempted by any other task. Usually this is the task number 1.

It recursively computes the value W until reaching $W_i^{n+1} = W_i^n$, which implies the solution is in a *steady-state* and the response time of task i is $R_i = W_i^n$.

It can also finish when the **release time** is greater than the **deadline** for any task.

Critical Time

Critical time refers to instants where a lower priority task is activated at the same time as a higher priority one.

Absolute Deadline

An *absolute deadline* is defined as:

$$d_i = \phi_i + kT_i + D_i$$

The value ϕ_i refers to the initial phase.

Contribution

The contribution of a task refers to how much time is uses of an interval. It is denoted as $\eta_i(t_a, t_b)$.

Processor on Demand Criterion

It checks that at any interval, the required computation by a task set is not greater than the available time. Using [response time](#) r and [absolute deadline](#) d , it is defined as:

$$g(t_1, t_2) = \sum_{\substack{r_{i,k} \geq t_1 \\ d_{i,k} \leq t_2}} C_i$$

A set of tasks are schedulable if in any time interval the processor demand does not exceed the available time.

$$\forall t_1, t_2 : g(t_1, t_2) \leq t_2 - t_1$$

To test this, it is required to find the instances where the task is contributing between t_1 and t_2 . Also taking into account the [contribution](#):

$$g(t_1, t_2) = \sum_{i=1}^n \eta_i(t_1, t_2) C_i = \sum_{i=1}^n \max(0, \lfloor \frac{t_2 + T_i - D_i}{T_i} \rfloor - \lceil \frac{t_1}{T_i} \rceil) C_i$$

Given that:

$$\begin{cases} \forall i & : \phi_i = 0 \\ t_1 & = 0 \\ t_2 & = L \end{cases}$$

Schedulability is **ensured** if $g(0, L) \leq L$. The formula simplifies to:

$$g(0, L) = \sum_{i=1}^n \eta_i(0, L) C_i = \sum_{i=1}^n \lfloor \frac{L + T_i - D_i}{T_i} \rfloor C_i$$

If $D_i = T_i$ then it is simplified to:

$$g(0, L) = \sum_{i=1}^n \lfloor \frac{L}{T_i} \rfloor C_i$$

Arrival Time

Tasks might *arrive* at a different time than when its period starts. It is ensured that its arrival time is fitted:

$$a_i + c_i \leq d_i$$

Aperiodic tasks

An *aperiodic task* is a task that has an [arrival time](#) other than 0.

Cyclic Scheduler

A cyclic scheduler is one where the scheduler chooses which task will be executed. This scheduling is composed of 2 elements:

- **Algorithm:** It determines how the tasks are executed
- **Analysis:** It guarantees timing constraints

In these systems, the time constraints are ensured by design. It needs a **time mechanism** that triggers the scheduler periodically.

Cyclic Scheduler: Requisites

The baseline of these systems are defined as:

- 1 processor.
- Static tasks.
- Periodic tasks.
- No *precedence* among tasks.
- The **WCET** is known for all tasks. They are all less or equal to their **deadline**.
- **Deadlines** of each task are *equal* to their **period**.

Cyclic Scheduler: Methodology

For a *cyclic scheduler* to be schedulable it needs to meet 2 criteria:

- The **utilization factor** must be *less or equal* than 1
- Using the **hyperperiod** find the **secondary period** that satisfies all of its conditions.

Cyclic Scheduler: Pros

- It is **static**, simple, easy to handle and robust
- **Deadlines** are ensured by design
- No concurrency nor **preemption**
- No mutual exclusion
- Low-Level scheduler

Cyclic Scheduler: Cons

- Not flexible
- Segmentation of tasks increases the complexity
- Not adequate for *sporadic tasks*
- Hard to find task allocation within few frames
- Low-Level scheduler

Rate Monotonic

Since each task has different rates ([periods](#)) we can order them by that criteria. High frequency tasks have higher priorities. Formally written:

$$\forall \tau_i, \tau_j : T_i < T_j \Rightarrow P_i > P_j$$
$$P_i \propto \frac{1}{T_i}$$

This means that the priorities are proportional to the inverse of the [period](#). The system will run the task with highest priority on each tick. This implies that [preemption](#) is possible.

Other important properties:

- The schedulability analysis attempts to know in advance the if all [release times](#) occurs before its [deadline](#)- The analysis is only performed at the [critical time](#).
- At runtime, the scheduler checks at each tick, the task with highest priority and dispatches it.

Rate Monotonic: Requisites

These systems have the same [base requisites](#) as a [cyclic scheduler](#). It adds 2 new requirements:

- 1 processor.
- Static tasks.
- Periodic tasks.
- No *precedence* among tasks.
- The [WCET](#) is know for all tasks. They are all less of equal to their [deadline](#).
- [Deadlines](#) of each task are *equal* to their [period](#).
- **Tasks can be preempted**
- **Real Time kernels uses *fixed priorities***

Rate Monotonic: Methodology

For a *rate monotonic scheduler* to work it needs to fulfill one of 2 *sufficient* conditions:

- [Factor utilization factor sum](#)
- [Hyperbolic bound](#)

If neither of these 2 conditions are satisfied, check [response time analysis](#) for the critical section.

Rate Monotonic: Methodology > Condition 1

This first condition requires the sum of the [utilization factor](#) of all tasks needs to be *less than or equal* to $n(\sqrt[n]{2} - 1)$. Formally written:

$$U_{total} = \sum_{i=1}^n U_i \leq n(\sqrt[n]{2} - 1)$$

Rate Monotonic: Methodology > Condition 2

This condition checks the following formula:

$$\prod_{i=1}^n (U_i + 1) \leq 2$$

Rate Monotonic: Pros

- At design time, priorities are configured based on rate of occurrence of each task
- 2 sufficient conditions
- [Response time analysis](#) being a necessary and sufficient condition
- Feasibility of scheduler is checked at critical time (less work to do)

Rate Monotonic: Cons

- It allows [preemption](#)
- Depends a lot on CPU utilization.

$$\lim_{n \rightarrow \infty} U_{\text{total}} = \lim_{n \rightarrow \infty} n(\sqrt[n]{2} - 1) = \ln(2) = 0.6931\dots$$

Deadline Monotonic

Deadline monotonic is a variant of [rate monotonic](#), used for tasks with [deadlines less or equal](#) than their [period](#).

$$\forall \tau_i, \tau_j : D_i < D_j \Rightarrow P_i > P_j$$
$$P_i \propto \frac{1}{D_i}$$

This means that the priorities are proportional to the inverse of the [deadline](#). The system will run the task with highest priority on each tick. This implies that [preemption](#) is possible.

Other important properties:

- The schedulability analysis attempts to know in advance the if all [release times](#) occurs before its [deadline](#)
- The analysis is only performed at the [critical time](#).
- At runtime, the scheduler checks at each tick, the task with highest priority and dispatches it.

Deadline Monotonic: Requisites

- 1 processor.
- Static tasks.
- Periodic tasks.
- No *precedence* among tasks.
- The [WCET](#) is know for all tasks. They are all less of equal to their [deadline](#).
- [Deadlines](#) of each task **are less or equal** to their [period](#).
- **Tasks can be preempted**
- **Real Time kernels uses *fixed priorities***

Deadline Monotonic: Methodology

It only applies [response time analysis](#), not using the 2 other sufficient conditions of [rate monotonic](#).

Deadline Monotonic: Pros

- At design time, priorities are configured based on rate of occurrence of each task
- [Response time analysis](#) being a necessary and sufficient condition

Deadline Monotonic: Cons

- Preemption
- Performance depends on system ticks

Earliest Deadline First

At run time, each system tick will check the priorities of active tasks and priorities are modified if needed.

$$\forall t, \tau_i, \tau_j : d_i < d_j \Rightarrow P_i > P_j$$
$$P_i \propto \frac{1}{d_i}$$

Where d is the [absolute deadline](#). This means that priorities are proportional to the inverse of the [absolute deadline](#). At each system tick, the scheduler looks for the highest priority tasks among active ones. This implies that [preemption](#) is possible.

Other important properties:

- The schedulability analysis attempts to know in advance the if all [release times](#) occurs before its [deadline](#)

Earliest Deadline First: Requisites

- 1 processor
- ~~Static tasks~~
- Periodic tasks
- ~~No precedence among tasks~~
- The [WCET](#) is know for all tasks. They are all less of equal to their [deadline](#)
- [Deadlines](#) of each task **are less or equal** to their [period](#)
- **Periodic or aperiodic tasks**
- **It allows precedence among tasks**
- **Tasks can be preempted**
- **Real Time kernels uses *dynamic priorities***

Earliest Deadline First: Methodology

There are 2 cases to check:

1. If $D_i = T_i$ then the necessary and sufficient condition is $U_{total} \leq 1$
2. If $D_i < T_i$ then it is necessary to check [processor on demand criterion](#)

Earliest Deadline First: Methodology > [Processor on Demand Criterion](#)

To check [processor on demand criterion](#) it is required to find L :

$$L \geq \sum_{i=1}^n (\lfloor \frac{L - D_i}{T_i} \rfloor + 1) C_i$$

Using this value, find $g(0, L) \leq L$ with the [absolute deadlines](#) obtained from the following set:

$$\begin{cases} D &= [d_k | d_k \leq \min(H, L^*)] \\ H &= \text{lcm}(T_1, T_2, \dots, T_n) \\ L^* &= \frac{\sum_{i=1}^n (T_i - D_i) U_i}{1 - U_{total}} \end{cases}$$

Earliest Deadline First: Precedence

One special case of [EDF](#) consists of task having precedence, meaning $\tau_a \rightarrow \tau_b$. The main effect is that both [arrival time](#) and [deadlines](#) are modified::

$$\tau_a \rightarrow \tau_b : \begin{cases} a_b^* &= a_a + C_a \\ d_a^* &= d_b + C_b \end{cases}$$

Where a is the arrival time

Given that a set J of tasks can transform into a set J^* of tasks, both the [arrival times](#) and [deadlines](#) must be modified:

From the set of tasks J to J^*

Arrival time :

1. Select task τ_i with a modified immediate predecessor
2. $\max_{\text{arrival}} = \max_{\tau_k \rightarrow \tau_i} (a_k^* + C_k), k \in J$
3. Set $a_i^* = \max\{a_i, \max_{\text{arrival}}\}$

Deadline :

1. Select task τ_i with a modified immediate successor
2. $\min_{\text{deadline}} = \min_{\tau_k \rightarrow \tau_i} (d_k^* - C_k), k \in J$
3. Set $d_i^* = \min\{d_i, \min_{\text{deadline}}\}$

Earliest Deadline First: New Task

Given that [EDF](#) checks at runtime, you can add tasks to it. This will transform the set of tasks J to J_{new} .

For this to work, the following must be true:

With $J_{\text{new}} \cup J = J'$:

for all i in J' : $f_i \leq d_i$

$$\begin{cases} f_i &= \sum_{k=1}^n b_k \\ b_k &: \text{remaining } \textit{Worst Case Execution Time} \text{ of a task } i \text{ in } J' \end{cases}$$

Basically check if adding this new task, with the remaining times up to the extra time for not reaching the **deadline** of each tasks, it will fit.

Earliest Deadline First: Pros

- Uses *dynamic priorities* set at runtime
- If $T_i = D_i$, then $U_{total} \leq 1$ is a sufficient condition
- It accepts **aperiodic tasks**

Earliest Deadline First: Cons

- It is less predictable and less controllable, when trying to reduce response time
- It requires more overhead
- Overload can lead to a domino effect